

# Diabetes Prediction

```
In [1]: import numpy as np
import pandas as pd
import pandera.pandas as pa
import altair as alt
import os
import warnings
from sklearn.preprocessing import StandardScaler
from sklearn.compose import make_column_transformer, make_column_selector
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import BernoulliNB
from sklearn.pipeline import make_pipeline
from ucimlrepo import fetch_ucirepo
from sklearn.metrics import (
    accuracy_score, make_scorer, fbeta_score,
    recall_score, precision_score, ConfusionMatrixDisplay
)
from deepchecks.tabular.checks import FeatureLabelCorrelation, FeatureFeature
from deepchecks.tabular import Dataset

alt.data_transformers.disable_max_rows()
warnings.filterwarnings("ignore", category=FutureWarning, module="deepchecks")
np.random.seed(522)
```

```
/opt/conda/lib/python3.11/site-packages/deepchecks/core/serialization/dataframe/html.py:16: UserWarning:
```

```
pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
```

## Summary

In this project we attempt to build a model to predict diabetes disease. We compared a decision tree model and naive bayes model and found the decision tree is stronger in this context. We used f2-score as our scoring function because detecting diabetes is the priority: a false negative could be much worse than a false positive.

In the test dataset: the decision tree model correctly detected 8283 of 10604 positive cases (recall rate is about 78%). This result does come at a fairly significant cost in terms of false positives (precision rate is about 30%) with 19650 false positives. Depending on the actual cost of false positive this may need significant improvement to be a viable screening model.

# Introduction

In Canada and the USA approximately 10% of people are living with diabetes. In Canada in 2023 approximately 3.7 million people were living with diabetes and in the USA in 2021 approximately 38.4 million people were living with diabetes. In the USA it is the 8th leading cause of death. Globally an estimated 44% of people living with diabetes are undiagnosed. (Snapshot of Diabetes in Canada, 2023; Rios et al., 2017; Stafford et al., 2025)

In this project we try to predict diabetes disease based on common health factors. A reliable model could help to prescreen people and recommend following up with a physician for people who are at risk. Given the large number of people living with undiagnosed diabetes this could potentially have a significant positive impact on world health.

The analysis uses the American CDC Behavioural Risk Factor Surveillance System (BRFSS) 2015 Diabetes Health Indicators dataset (UCI ID 891), containing 253,680 survey responses with 21 health-related features and a binary diabetes outcome (0 = no diabetes/pre-diabetes, 1 = diabetes).

No missing values were present and all features were already encoded numerically. The target classes are imbalanced ( $\approx 86\%$  non-diabetic,  $\approx 14\%$  diabetic).

## Methods and Results

The analysis uses the CDC Behavioural Risk Factor Surveillance System (BRFSS) 2015 Diabetes Health Indicators dataset (UCI ID 891), containing 253,680 survey responses with 21 health-related features and a binary diabetes outcome (0 = no diabetes/pre-diabetes, 1 = diabetes). (Dane and Teboul, 2021) No missing values were present and all features were already encoded numerically. The target classes are heavily imbalanced ( $\approx 86\%$  non-diabetic,  $\approx 14\%$  diabetic).

## EDA

Group-wise mean differences revealed the strongest risk factors for diabetes:

- PhysHlth (days of poor physical health)
- BMI
- Age
- MentHlth (days of poor mental health)
- GenHlth (self-rated general health)

Weakest factors

- HvyAlcoholConsump
- Fruits
- Veggies
- PhysActivity
- Education
- Income

Box plots of the top five predictors clearly separate the diabetic and non-diabetic groups.

## Modeling Approach

The data were split 70/30 into training and test sets with stratification on the target. Two classifiers were trained and tuned using 5-fold cross-validated grid search with **f2-score** as the scoring metric. We chose to use f2-score because it is more appropriate than accuracy or f1 because we don't want to miss true positives.

### 1. Decision Tree (class\_weight='balanced')

Hyperparameters: max\_depth ∈ {6,8,10,12,14}, min\_samples\_leaf ∈ {175, 200, 225, 250}

**Best parameters:** max\_depth=10, min\_samples\_leaf=200

**Best CV f2-score** = 0.5908

### 2. Bernoulli Naive Bayes (with StandardScaler preprocessing)

Hyperparameters: alpha ∈ {1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2, 1e3, 1e4}

**Best parameters:** alpha=1e-3

**Best CV f2-score** = 0.4453

## Results

Model	Test Accuracy	Test f2-score	Test recall	Test precision
Decision Tree	0.706	<b>0.587</b>	<b>0.783</b>	0.293
Naive Bayes	<b>0.814</b>	0.460	0.489	<b>0.373</b>

Table: 1

## Load Data

```
In [2]: # fetch dataset
cdc_diabetes_health_indicators = fetch_uci_repo(id=891)

# data (as pandas dataframes)
X = cdc_diabetes_health_indicators.data.features
y = cdc_diabetes_health_indicators.data.targets
```

# Validate Data Before Saving

## X Verification

In [3]:

```
#####
# Checks verified:
# Correct data file format
# Correct column names
# No empty observations (and by extension Missingness not beyond expected thresholds)
# Correct data types in each column
# No outlier or anomalous values

Maximum allowable ranges for the numeric_features were determined from
the schema description in https://archive.ics.uci.edu/dataset/891/cdc+diabetes
#####

# Check that X is a Pandas DataFrame
if not isinstance(X, pd.DataFrame):
    raise TypeError("X object obtained is not a Pandas Dataframe.")

# Create a list of expected column names
column_names = ["HighBP", "HighChol", "CholCheck", "BMI", "Smoker", "Stroke", "HeartDiseaseOrAttack",
                "Fruits", "Veggies", "HvyAlcoholConsump", "AnyHealthcare",
                "NoDocbcCost", "GenHlth", "MentHlth", "PhysHlth", "DiffWalk", "Sex"]

# Create a list of column names that are binary features
binary_features = ["HighBP", "HighChol", "CholCheck", "Smoker", "Stroke", "HeartDiseaseOrAttack",
                    "Fruits", "Veggies", "HvyAlcoholConsump", "AnyHealthcare",
                    "NoDocbcCost", "DiffWalk", "Sex"]

# Define allowable ranges for numeric features
numeric_features = {"BMI": (5, 256),
                     "GenHlth": (1, 5),
                     "MentHlth": (0, 30),
                     "PhysHlth": (0, 30),
                     "Age": (1, 13),
                     "Education": (1, 6),
                     "Income": (1, 8)}

# Check that all expected columns are present
schema_dict = {}
for col_name in binary_features:
    schema_dict[col_name] = pa.Column(int, pa.Check.between(0, 1), nullable = False)

# Add numeric features to schema with their respective ranges
for col_name in numeric_features.keys():
    schema_dict[col_name] = pa.Column(int, pa.Check.between(numeric_features[col_name][0], numeric_features[col_name][1]), nullable = False)

schema = pa.DataFrameSchema(schema_dict)

schema.validate(X, lazy = True)
```

Out[3]:

	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorAttack	
0	1	1	1	40	1	0		0
1	0	0	0	25	1	0		0
2	1	1	1	28	0	0		0
3	1	0	1	27	0	0		0
4	1	1	1	24	0	0		0
...	...	...	...	...	...	...	...	...
253675	1	1	1	45	0	0		0
253676	1	1	1	18	0	0		0
253677	0	0	1	28	0	0		0
253678	1	0	1	23	0	0		0
253679	1	1	1	25	0	0		1

253680 rows × 21 columns

## Y Verification

In [4]:

```
# Use same template as X verification, only one column
if not isinstance(y, pd.DataFrame):
    raise TypeError("y object obtained is not a Pandas Dataframe.")

schema_dict = {}

target_name = ["Diabetes_binary"]

for col_name in target_name:
    schema_dict[col_name] = pa.Column(int, pa.Check.between(0,1), nullable =)

schema = pa.DataFrameSchema(schema_dict)

schema.validate(y, lazy = True)
```

Out [4] :

Diabetes_binary	
0	0
1	0
2	0
3	0
4	0
...	...
253675	0
253676	1
253677	0
253678	0
253679	1

253680 rows × 1 columns

## Save Raw Data

In [5] :

```
# check if raw folder exists
raw_data_path = "../data/raw"

if not os.path.exists(raw_data_path):
    os.makedirs(raw_data_path)

## Save Raw Data
X.to_csv("../data/raw/diabetes_raw_features.csv")
y.to_csv("../data/raw/diabetes_raw_targets.csv")
```

## Data Wrangling

In [6] :

```
# No major cleaning needed – dataset is already very clean!
# Combine features and targets to get a overview of the full data set
df = X.copy()
df['diabetes'] = y

# Quick info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 253680 entries, 0 to 253679
Data columns (total 22 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   HighBP          253680 non-null   int64  
 1   HighChol        253680 non-null   int64  
 2   CholCheck       253680 non-null   int64  
 3   BMI             253680 non-null   int64  
 4   Smoker          253680 non-null   int64  
 5   Stroke          253680 non-null   int64  
 6   HeartDiseaseorAttack 253680 non-null   int64  
 7   PhysActivity    253680 non-null   int64  
 8   Fruits          253680 non-null   int64  
 9   Veggies         253680 non-null   int64  
 10  HvyAlcoholConsump 253680 non-null   int64  
 11  AnyHealthcare   253680 non-null   int64  
 12  NoDocbcCost    253680 non-null   int64  
 13  GenHlth         253680 non-null   int64  
 14  MentHlth        253680 non-null   int64  
 15  PhysHlth        253680 non-null   int64  
 16  DiffWalk        253680 non-null   int64  
 17  Sex              253680 non-null   int64  
 18  Age              253680 non-null   int64  
 19  Education        253680 non-null   int64  
 20  Income           253680 non-null   int64  
 21  diabetes         253680 non-null   int64  
dtypes: int64(22)
memory usage: 42.6 MB
```

In [7]: `df.head()`

Out[7]:

	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorAttack	PhysA
0	1	1	1	40	1	0	0	0
1	0	0	0	25	1	0	0	0
2	1	1	1	28	0	0	0	0
3	1	0	1	27	0	0	0	0
4	1	1	1	24	0	0	0	0

5 rows × 22 columns

In [8]: `df.tail()`

Out[8]:

	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorAttack
253675	1	1	1	45	0	0	0
253676	1	1	1	18	0	0	0
253677	0	0	1	28	0	0	0
253678	1	0	1	23	0	0	0
253679	1	1	1	25	0	0	1

5 rows × 22 columns

## Data Summary

All the features in this dataset were selected by clinician to be relevant to a person having diabetes. We will use all features in our model.

In [9]:

df.describe()

Out[9]:

	HighBP	HighChol	CholCheck	BMI	Smoker
count	253680.000000	253680.000000	253680.000000	253680.000000	253680.000000
mean	0.429001	0.424121	0.962670	28.382364	0.44316
std	0.494934	0.494210	0.189571	6.608694	0.49676
min	0.000000	0.000000	0.000000	12.000000	0.00000
25%	0.000000	0.000000	1.000000	24.000000	0.00000
50%	0.000000	0.000000	1.000000	27.000000	0.00000
75%	1.000000	1.000000	1.000000	31.000000	1.00000
max	1.000000	1.000000	1.000000	98.000000	1.00000

8 rows × 22 columns

In [10]:

```
"""
These checks ensure data quality by verifying:
- Checking if duplicate rows exist in the dataset, preventing redundant data
  (We are not dropping duplicates automatically to avoid unintentional data loss)
  (For the purpose of this dataset, because we already finished the analysis)
- Checking for outlier in ("BMI", "MentHlth", "PhysHlth") columns to see if
  which helps maintain data integrity and model robustness.
  (For the purpose of this dataset, we will accept the outlier value because it's
  - Categorical validation is not applicable for this dataset
"""

def iqr_outliers(series: pd.Series) -> bool:
    """Return True if all values are within 1.5 * IQR (no extreme outliers)"""
    Q1 = series.quantile(0.25)
```

```

Q3 = series.quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
return series.between(lower_bound, upper_bound).all()

# Define the schema with meaningful checks
diabetes_schema = pa.DataFrameSchema(
    columns={
        "BMI": pa.Column(float, nullable=False),
        "MentHlth": pa.Column(float, nullable=False),
        "PhysHlth": pa.Column(float, nullable=False),
    },
    checks=[
        # 1. No duplicate rows
        pa.Check(
            lambda df: ~df.duplicated().any(),
            error="DUPLICATE_ROWS: Found duplicate observations. Use df.drop_duplicates()",
        ),

        # 2. No extreme outliers in continuous variables using IQR rule
        pa.Check(
            lambda df: iqr_outliers(df["BMI"]),
            error="OUTLIERS_IN_BMI: Extreme outliers detected in BMI (beyond IQR range)",
        ),
        pa.Check(
            lambda df: iqr_outliers(df["MentHlth"]),
            error="OUTLIERS_IN_MENTHLTH: Extreme values in MentHlth (beyond IQR range)",
        ),
        pa.Check(
            lambda df: iqr_outliers(df["PhysHlth"]),
            error="OUTLIERS_IN_PHYSHLTH: Extreme values in PhysHlth (beyond IQR range)",
        )
    ]
)

# Validate with lazy=True to see all errors at once
try:
    diabetes_schema.validate(df, lazy=True)
    print("All checks passed! Dataset is clean and ready for modeling.")
except pa.errors.SchemaErrors as e:
    print("Validation failed! See errors below:")
    print(e.failure_cases) # Shows detailed failure report

```

```

Validation failed! See errors below:
      column failure_case index    schema_context \
3        None         False  None  DataFrameSchema
4        None         False  None  DataFrameSchema
5        None         False  None  DataFrameSchema
6        None         False  None  DataFrameSchema
0       BMI        int64  None       Column
1  MentHlth        int64  None       Column
2  PhysHlth        int64  None       Column

                                         check  check_number
3 DUPLICATE_ROWS: Found duplicate observations. ...          0.0
4 OUTLIERS_IN_BMI: Extreme outliers detected in ...        1.0
5 OUTLIERS_IN_MENTHLTH: Extreme values in MentHl...        2.0
6 OUTLIERS_IN_PHYSHLTH: Extreme values in PhysHl...        3.0
0                               dtype('float64')           NaN
1                               dtype('float64')           NaN
2                               dtype('float64')           NaN

```

## Visualizations

```
In [11]: diabetes_count = pd.DataFrame(df['diabetes'].value_counts()).reset_index()

alt.Chart(diabetes_count).mark_bar().encode(
    x=alt.X('diabetes:0', title='Has Diabetes'),
    y="count",
    color="diabetes:N"
).properties(title='Count of Diabetes vs Non-Diabetes Records in Dataset')
```

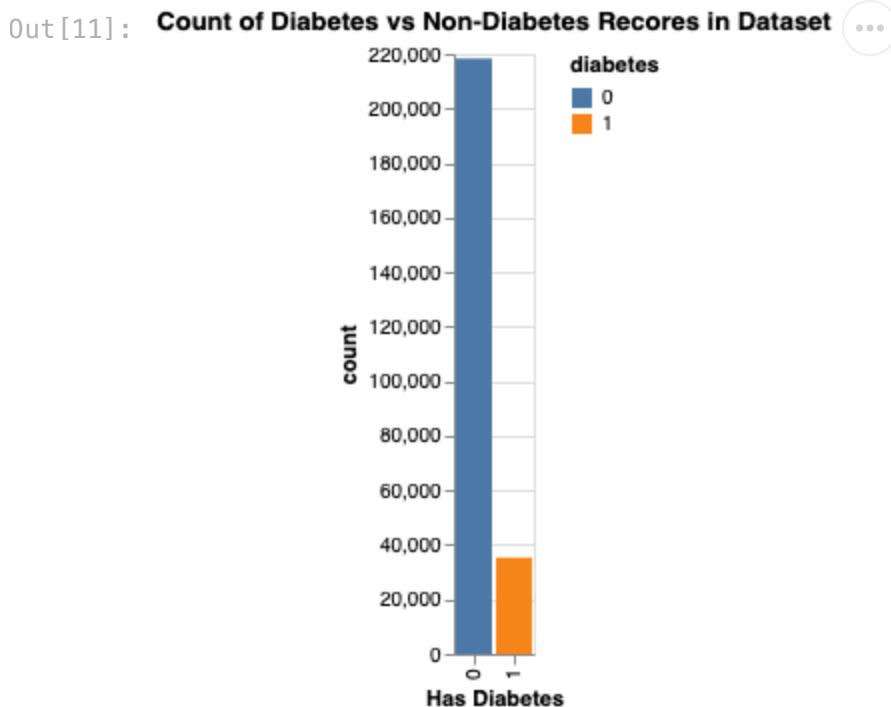


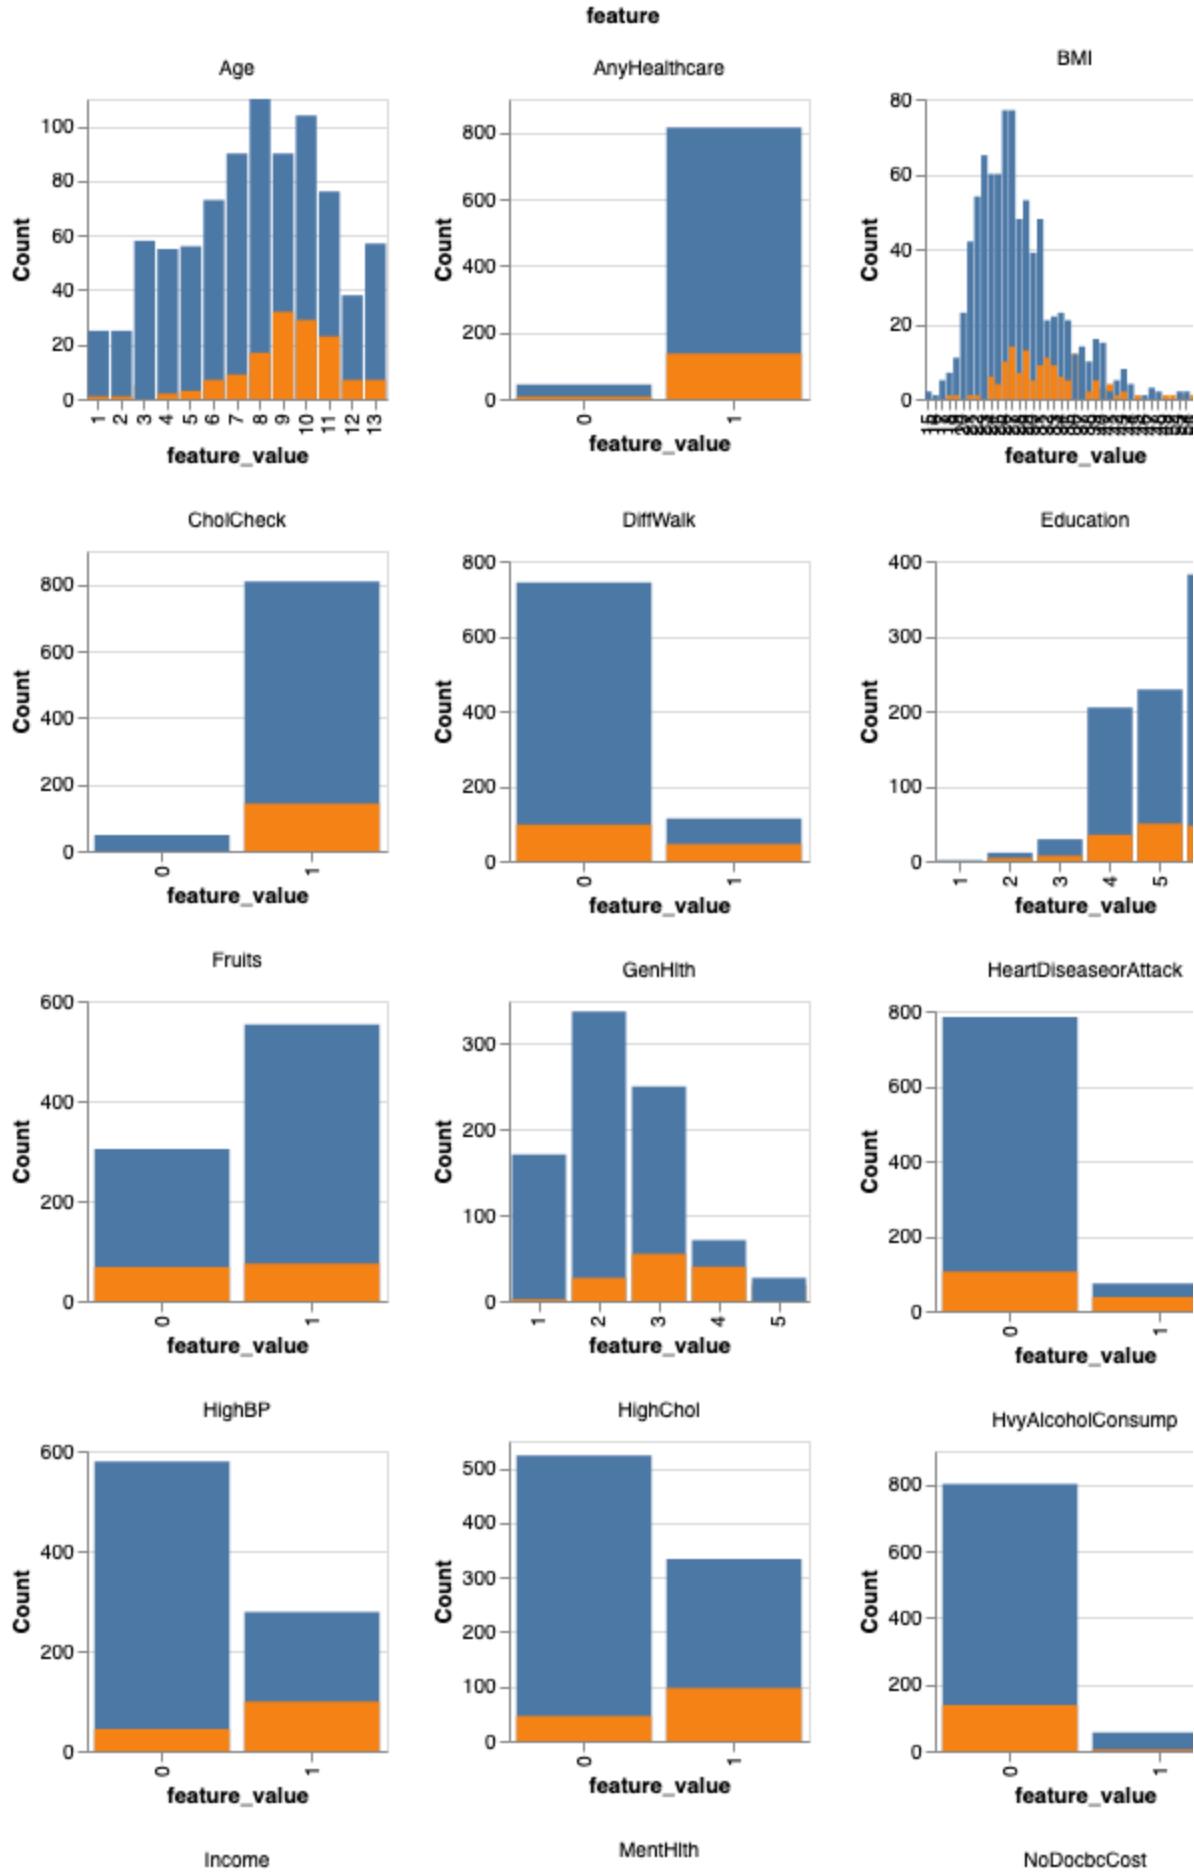
Figure 1

```
In [12]: df_sample = df.sample(n=1000, random_state=522)
features = df_sample.columns.to_list()
```

```
df_sample_long = pd.melt(df_sample, id_vars=["diabetes"], value_vars=feature)

histograms = alt.Chart(df_sample_long).mark_bar().encode(
    x=alt.X("feature_value:O", # Choose to use ordinal instead of quantitative
            title="Count"),
    y=alt.Y("count()", title="Count").stack(False),
    color=alt.Color("diabetes:N"),
).properties(
    width=150,
    height=150,
).facet(
    "feature:N",
    columns=3,
).resolve_scale(
    x="independent",
    y="independent",
).properties(
    title='Histograms of Features',
)
histograms
```

Out[12]: **Histograms of Features**



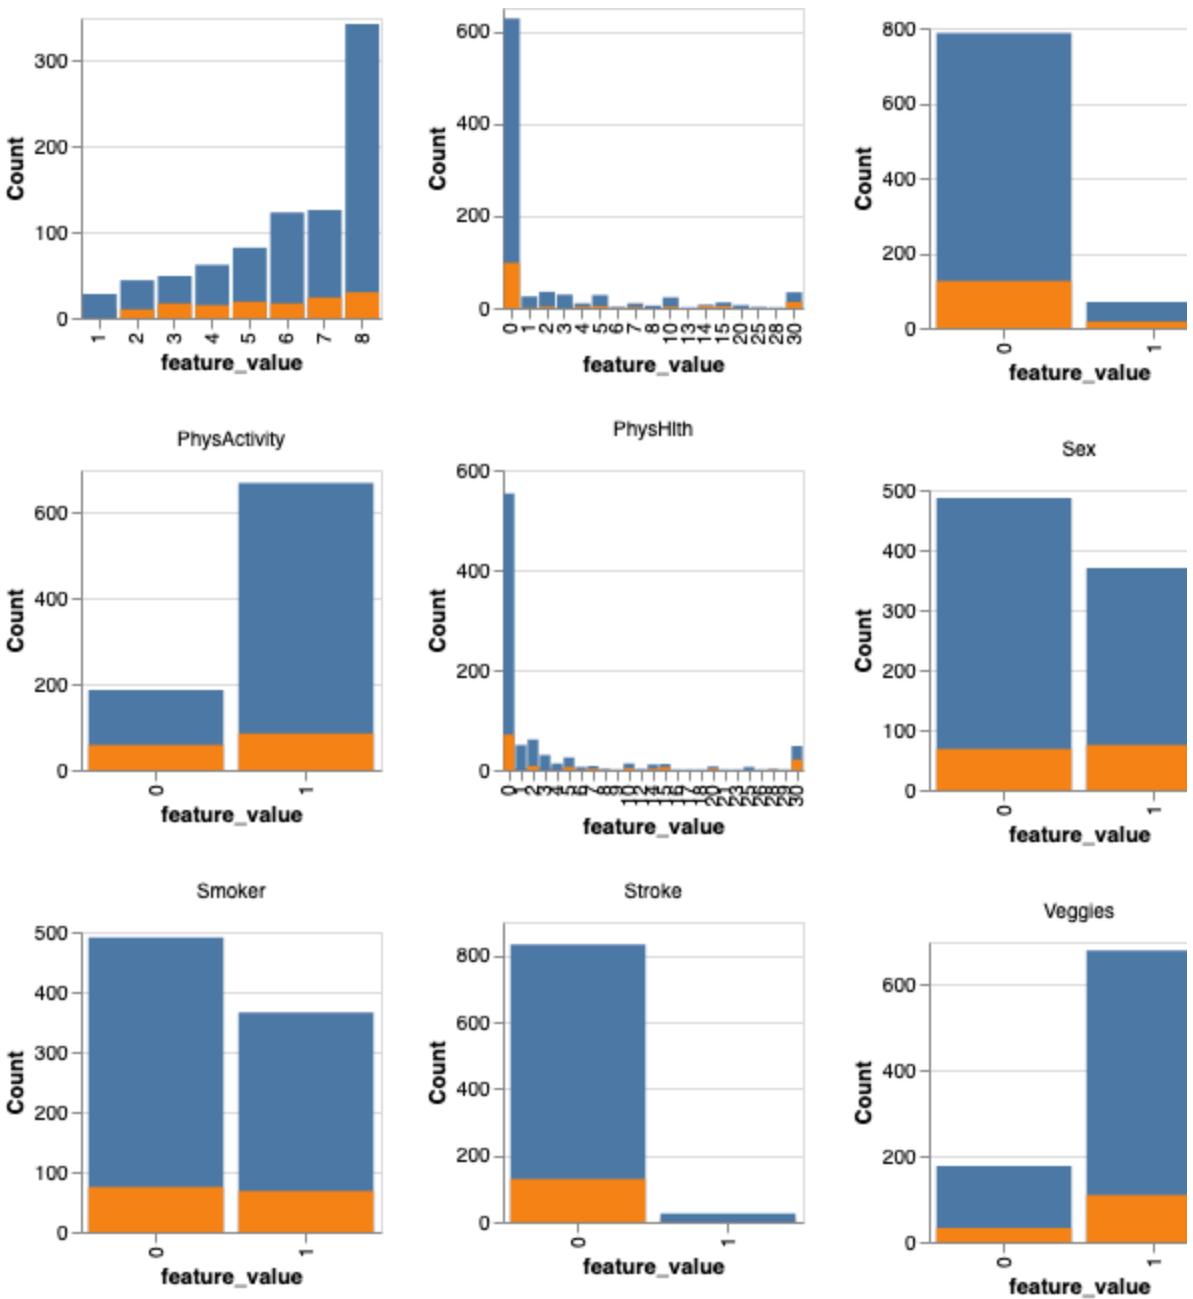


Figure 2

```
In [13]: # For binary features boxplot is not informative
non_binary_features = ['BMI', 'GenHlth', 'MentHlth', 'PhysHlth', 'Age', 'Ed'
df_sample_nonbinary = df_sample_long[df_sample_long["feature"].isin(non_bina
alt.Chart(df_sample_nonbinary).mark_boxplot().encode(
    x='diabetes:N',
    y='feature_value:Q',
    color='diabetes:N'
).facet(
    column='feature:N'
).properties(
    title='Boxplots for Non-Binary Features',
).resolve_scale(
```

```
y="independent"
)
```

Out [13]: **Boxplots for Non-Binary Features**

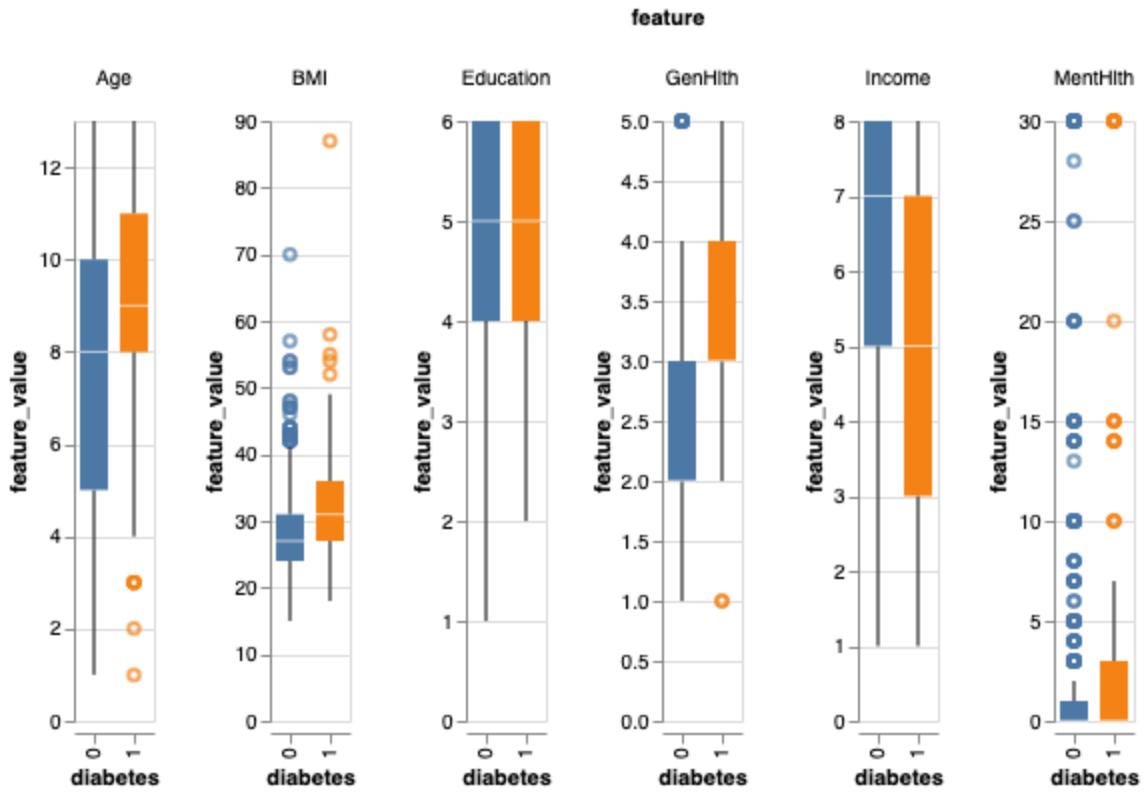


Figure 3

```
# split the data 70-30 split
train_df, test_df = train_test_split(
    df, test_size=0.3, random_state=522, stratify=df['diabetes']
)

# check if processed folder exists
processed_data_path = "../data/processed"

if not os.path.exists(processed_data_path):
    os.makedirs(processed_data_path)

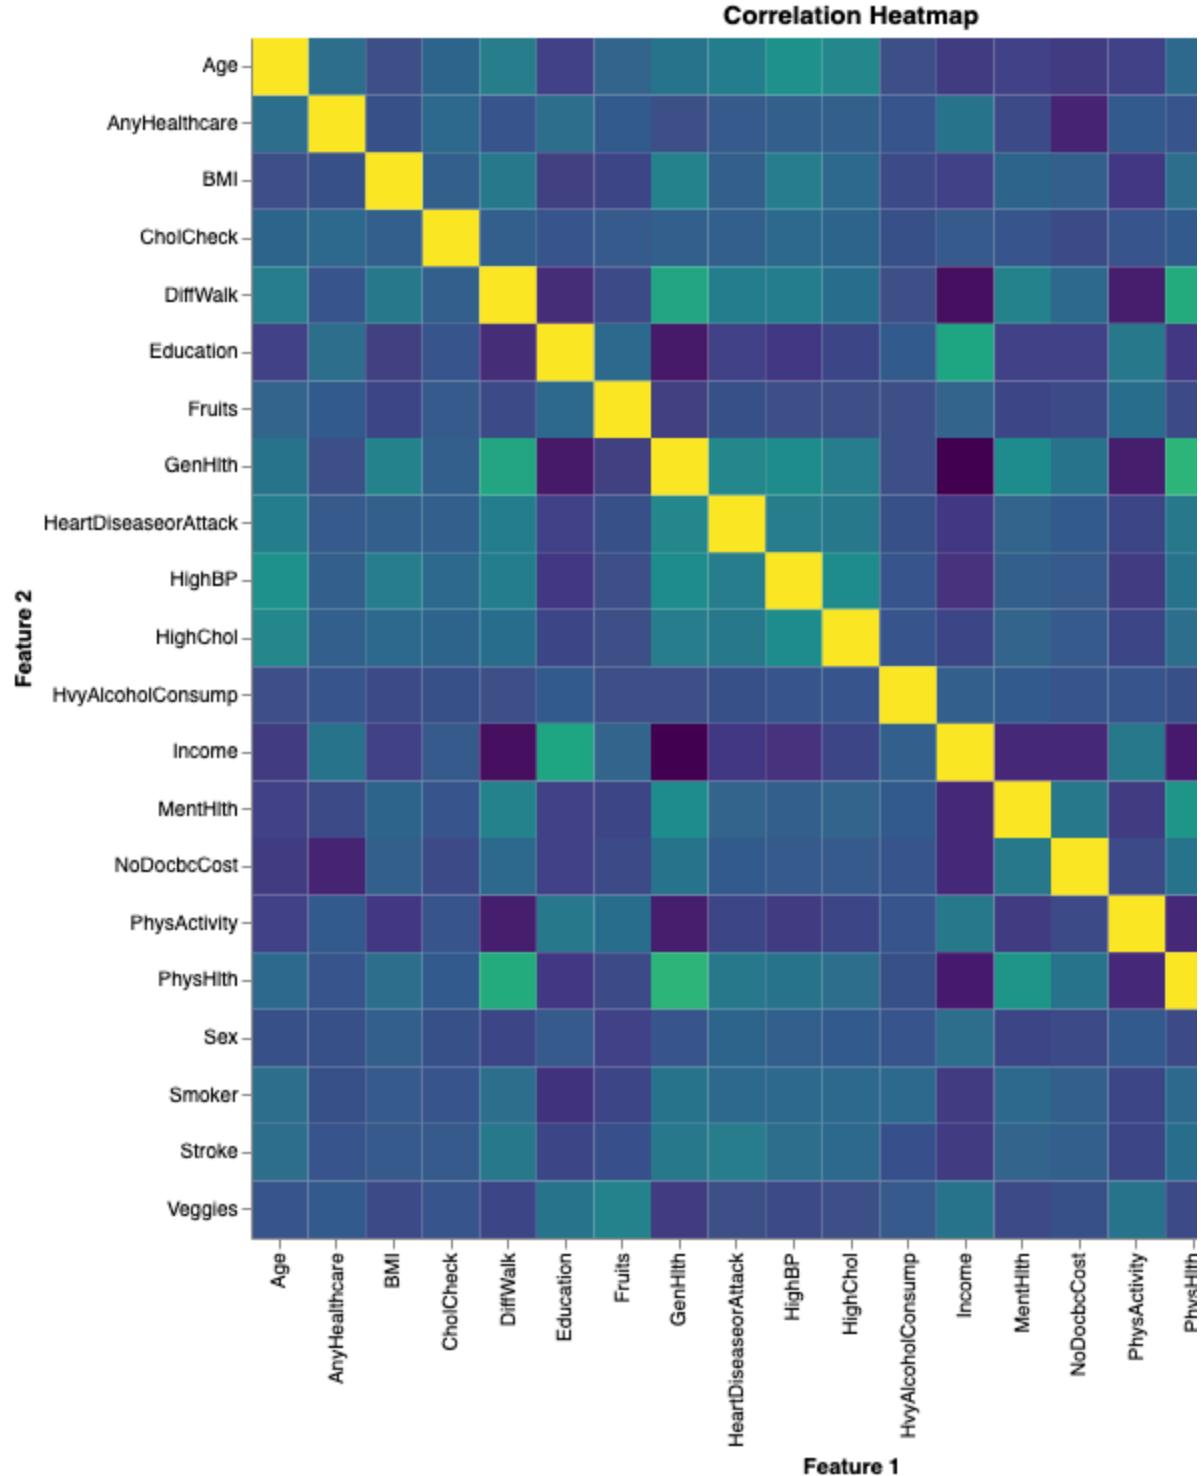
# Save processed data
train_df.to_csv(processed_data_path+"/diabetes_train.csv", index=False)
test_df.to_csv(processed_data_path+"/diabetes_test.csv", index=False)
```

```
X_train = train_df.drop('diabetes', axis=1)
y_train = train_df['diabetes']
X_test = test_df.drop('diabetes', axis=1)
y_test = test_df['diabetes']
```

```
# Create correlation chart
correlation_matrix = X_train.corr()
correlation_long = correlation_matrix.reset_index().melt(id_vars='index')
correlation_long.columns = ['Feature 1', 'Feature 2', 'Correlation']
```

```
alt.Chart(correlation_long).mark_rect().encode(
    x='Feature 1:0',
    y='Feature 2:0',
    color=alt.Color('Correlation:Q', scale=alt.Scale(scheme='viridis')),
    tooltip=['Feature 1', 'Feature 2', 'Correlation']
).properties(
    width=600,
    height=600,
    title="Correlation Heatmap"
)
```

Out[16]:



In [17]:

```
'''  
- Target/response variable follows expected distribution  
    (In the diabetes dataset, we expect the prevalence of diabetes to be around 10% to 15%)  
  
Validate training data for anomalous correlations:  
- Feature-label correlations (target vs features)  
- Feature-feature correlations (between features)  
  
Thresholds set based on domain knowledge.  
  
We perform these checks on training data only because including test data here  
and invalidate the evaluation of model generalization.  
'''  
  
# Target variable follows expected distribution  
# In this dataset: ~13–15% diabetes is normal, >30% or <5% is suspicious  
  
distribution_schema = pa.DataFrameSchema(  
    checks=[pa.Check(  
        lambda df: df.mean().between(0.05, 0.30),  
        error="ANOMALOUS_TARGET_DISTRIBUTION: Diabetes prevalence should be  
        between 0.05 and 0.30")  
    ]  
)  
  
distribution_schema.validate(train_df[['diabetes']], lazy=True)  
  
# Combine features and target into a dataset for validation  
diabetes_train_ds = Dataset(  
    df=train_df,  
    label="diabetes",  
    cat_features=[]  
)  
  
# Check that feature-label predictive power score (PPS) is below 0.9  
check_feat_lab_corr = FeatureLabelCorrelation().add_condition_feature_pps_le  
check_feat_lab_corr_result = check_feat_lab_corr.run(diabetes_train_ds)  
  
# Check that no feature pairs have correlation above 0.92  
check_feat_feat_corr = FeatureFeatureCorrelation().add_condition_max_number_<br/>  
    threshold=0.92, n_pairs=0  
)  
check_feat_feat_corr_result = check_feat_feat_corr.run(diabetes_train_ds)  
  
# Raise errors if any checks fail  
if not check_feat_lab_corr_result.passed_conditions():  
    raise ValueError("Feature-Label correlation exceeds the maximum acceptable  
    threshold of 0.9")  
  
if not check_feat_feat_corr_result.passed_conditions():  
    raise ValueError("Feature-feature correlation exceeds the maximum acceptab
```

## Classification Analysis

```
In [18]: f2_scorer = make_scorer(fbeta_score, beta=2)
```

```
In [19]: tree = DecisionTreeClassifier(random_state=522, class_weight='balanced')
```

```
tree_params = {
    'max_depth': [6, 8, 10, 12, 14],
    'min_samples_leaf': [175, 200, 225, 250]
}

tree_grid = GridSearchCV(tree, tree_params, cv=5, scoring=f2_scorer, n_jobs=-1)
tree_grid.fit(X_train, y_train)

best_tree = tree_grid.best_estimator_
print("Best Decision Tree params:", tree_grid.best_params_)
print("Best CV f2-score:", tree_grid.best_score_.round(4))
```

```
Best Decision Tree params: {'max_depth': 10, 'min_samples_leaf': 225}
Best CV f2-score: 0.5919
```

```
In [20]: preprocessor = make_column_transformer(
    (StandardScaler(), X_train.columns)
)
```

```
nb_pipe = make_pipeline(
    preprocessor,
    BernoulliNB()
)
```

```
nb_params = {'bernuullinb_alpha': [1e-3, 1e-2, 1e-1, 1e0, 1e1, 1e2, 1e3, 1e4]}
```

```
knn_grid = GridSearchCV(nb_pipe, nb_params, cv=5, scoring=f2_scorer, n_jobs=-1)
knn_grid.fit(X_train, y_train)
```

```
best_nb = knn_grid.best_estimator_
print("Best NB k:", knn_grid.best_params_)
print("Best CV f2-score:", knn_grid.best_score_.round(4))
```

```
Best NB k: {'bernuullinb_alpha': 0.001}
Best CV f2-score: 0.4604
```

```
In [21]: models = {
    'Decision Tree': best_tree,
    'Naive Bayes': best_nb
}
```

```
results = []
```

```
for name, model in models.items():
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)[:, 1]

    results.append({
        'Model': name,
        'Test Accuracy': round(accuracy_score(y_test, y_pred), 3),
        'Test f2-score': round(fbeta_score(y_test, y_pred, beta=2), 3),
        'Test recall': round(recall_score(y_test, y_pred), 3),
```

```

        'Test precision': round(precision_score(y_test, y_pred),3),
    })

score_df = pd.DataFrame(results)
score_df

```

Out[21]:

	Model	Test Accuracy	Test f2-score	Test recall	Test precision
0	Decision Tree	0.706	0.587	0.784	0.293
1	Naive Bayes	0.814	0.460	0.489	0.373

## Result Visualizations

In [22]:

```

score_melt = score_df.melt(id_vars='Model', var_name='Metric', value_name='Score')

alt.Chart(score_melt).mark_bar().encode(
    x='Model:N',
    y='Score:Q',
    color='Model:N',
    column='Metric:N'
).properties(
    title='Decision Tree vs Naive Bayes Performance on Test Set'
)

```

Out[22]: **Decision Tree vs Naive Bayes Performance on Test Set**

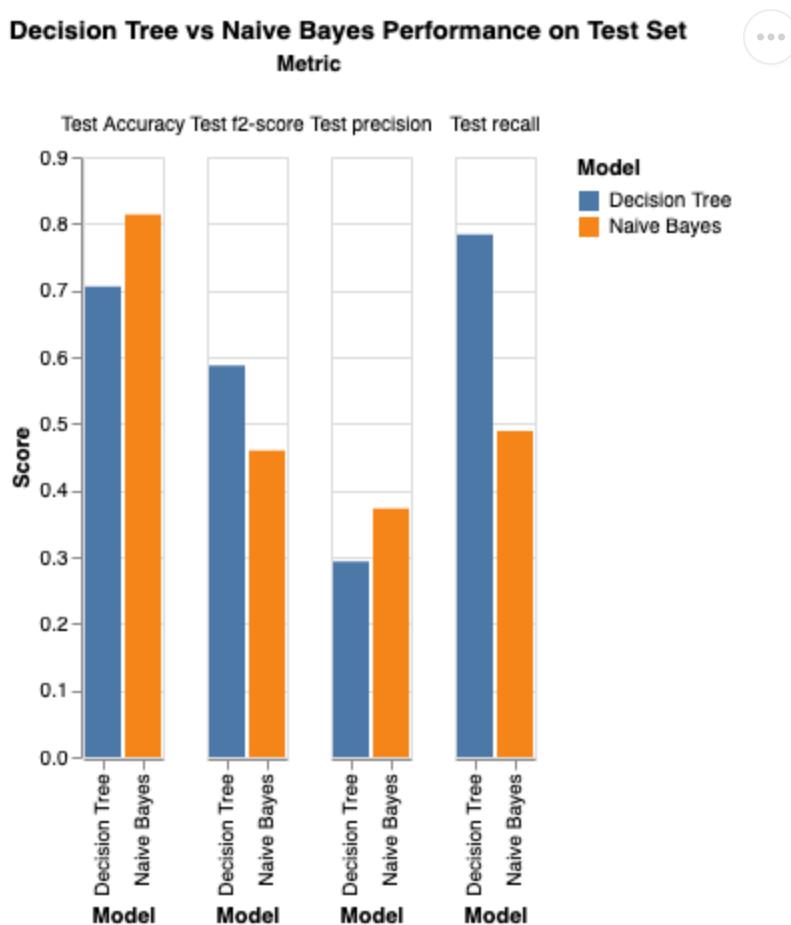


Figure 4

```
In [23]: # Confusion matrix for best model (decision tree)
```

```
ConfusionMatrixDisplay.from_estimator(  
    best_tree,  
    X_test,  
    y_test,  
    values_format="d",  
)
```

```
Out[23]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f3fddf  
15a50>
```

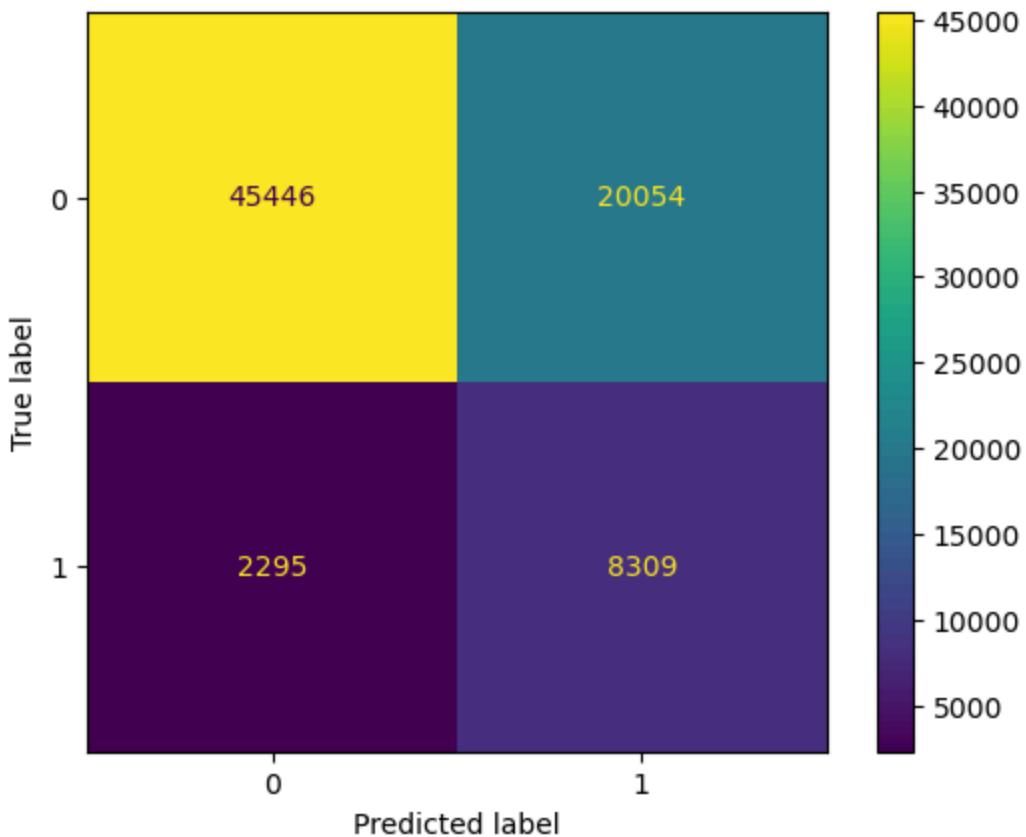


Figure 5

## Discussion

The current performance of the tree model is likely already good enough to offer some benefit in the real world given the large number of people with undiagnosed diabetes. However the recall score likely could be improved on and the precision score definitely leaves something to be desired.

We were surprised by the high rate of false positives. This might be an indication of how many non-diabetic people are at risk.

Further improvements to predicting diabetes could likely be found by 1) trying a wider variety of model type and using a wider hyperparameter search 2) possibly through

more feature engineering.

A future study could be done to find a smaller set of the most easy to obtain features. Such a model would be more usable by the average person. Some work is needed to determine this smaller number of easy to obtain features that doesn't significantly reduce model performance.

Another question is if a regression model could be made that predicts a persons risk as a percent chance of developing diabetes. Longitudinal data might be required for this type of prediction.

## References

Dane, Sohier, and Alex Teboul. Diabetes Health Indicators Dataset. 2021,  
<https://www.kaggle.com/datasets/alexteboul/diabetes-health-indicators-dataset/data>.

Kelly, Markelle, et al. The UCI Machine Learning Repository. 2021, Utilized the ucimlrepo library for data access. Further documentation on the library is located at  
<https://github.com/uci-ml-repo/ucimlrepo.archive.ics.uci.edu/ml>.

Python 3.12.12 documentation. 2021-2025,  
<https://docs.python.org/3.12/reference/index.html>.

Rios, Nilka Burrows, et al. Incidence of End-Stage Renal Disease Attributed to Diabetes Among Persons with Diagnosed Diabetes — United States and Puerto Rico. Morb Mortal Wkly Rep, 66(43), Nov. 2017, <http://dx.doi.org/10.15585/mmwr.mm6643a2>, pp. 1165–70.

Snapshot of Diabetes in Canada, 2023. 2023, <https://www.canada.ca/en/public-health/services/publications/diseases-conditions/snapshot-diabetes-canada-2023.html>.

Stafford, Laurn K, et al. "Global, regional, and national cascades of diabetes care, 2000–23: a systematic review and modelling analysis using findings from the Global Burden of Disease Study". The Lancet Diabetes & Endocrinology, 13(11), 2025, [https://doi.org/10.1016/S2213-8587\(25\)00217-7](https://doi.org/10.1016/S2213-8587(25)00217-7), pp. 924–34.